

Codify Your Data Infrastructure Using Terraform

Dewan Ahmed

Note: These slides have been modified from their original versions which were dedicated for a live audience.

THANKS TO ALL OUR SPONSORS!



Hi, I'm Dewan

- Senior Developer Advocate, Aiven
- New Brunswick, Canada
- Focus on app/data infrastructure
- Pro bono career coach



@aiven_io

in/diahmed | @DewanAhmed

Application Infrastructure



Why do we build applications?

To move **data**

Data Infrastructure

Streaming
Platform

Relational
Database

NoSQL
Database

Networking
(for data)

Monitoring
(for data)

Security
(for data)

Infrastructure-as-Code Principles

- Reproducibility
- Repeatability
- Disposability
- Consistency
- Ability to incorporate design changes

More in the book:
Infrastructure as Code (Kief Morris)

Infrastructure-as-Code Challenges

- Resistance to learning
- Configuration drift
- Duplicate work
- Security issues*
- Handling unicorns

*Source: <https://bridgecrew.io/blog/state-of-open-source-terraform-security-report-2020>

Handling Unicorns - Databases

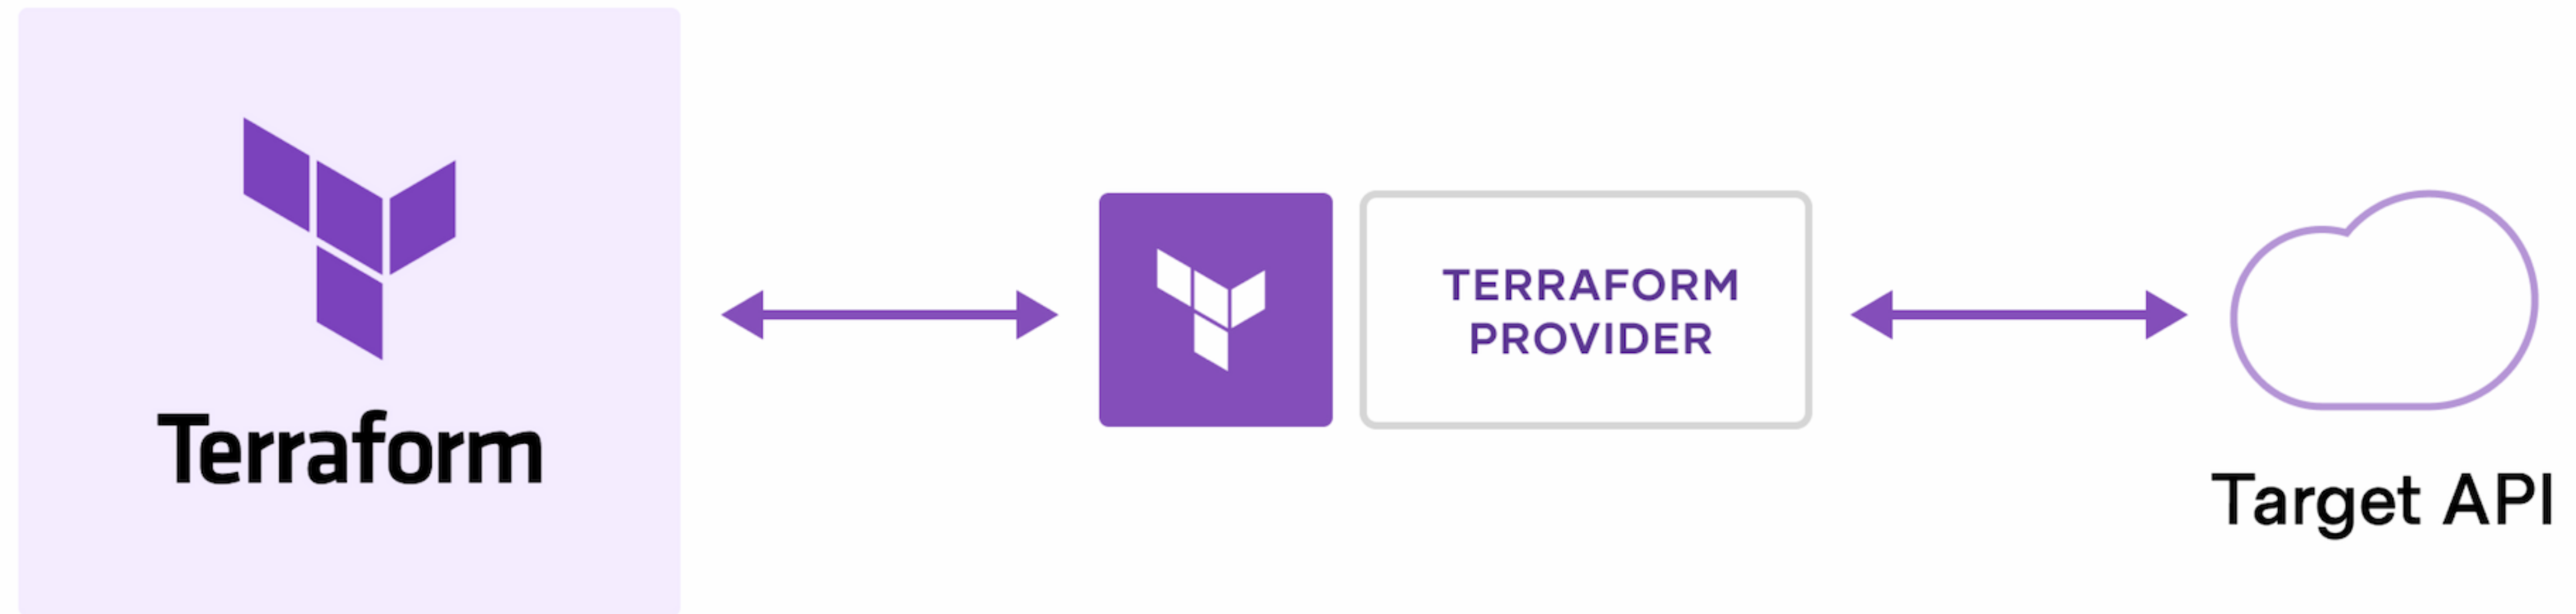
How to IaC without downtime

- Perform version upgrade
- Prevent version upgrade
- Change configurations

HashiCorp Terraform

- Open-source Infrastructure as Code tool
- Policy compliance and management
- Target cloud or on-prem resources
- Version-control
- Declarative

Aiven Terraform Provider



Source: <https://www.terraform.io/intro>

There was a demo in this slide.

<https://developer.aiven.io/docs/tools/terraform/get-started.html> to follow
the demo

Reproducibility

```
resource "aiven_m3aggregator" "demo-m3a" {  
  project      = var.project_name  
  cloud_name   = "google-northamerica-northeast1"  
  plan        = "business-8"  
  service_name = join("-", [var.service_name_prefix, "m3a"])  
  
  m3aggregator_user_config {  
    m3aggregator_version = 1.5  
  }  
}
```

```
# European Postgres Service
resource "aiven_pg" "dewan-avn-eu-pg" {
  project      = var.project_name
  cloud_name   = "aws-eu-west-2" # London
  plan        = "business-8"     # Primary + read replica
  service_name = join("-", [var.service_name_prefix, "postgres-eu"])
  termination_protection = true
}

# US Postgres Service
resource "aiven_pg" "dewan-avn-us-pg" {
  project      = var.project_name
  cloud_name   = "do-nyc"        # New York
  plan        = "business-8"     # Primary + read replica
  service_name = join("-", [var.service_name_prefix, "postgres-us"])
  termination_protection = true
}

# Asia Postgres Service
resource "aiven_pg" "dewan-avn-as-pg" {
  project      = var.project_name
  cloud_name   = "google-asia-southeast1" # Singapore
  plan        = "business-8"             # Primary + read replica
  service_name = join("-", [var.service_name_prefix, "postgres-asia"])
  termination_protection = true
}
```

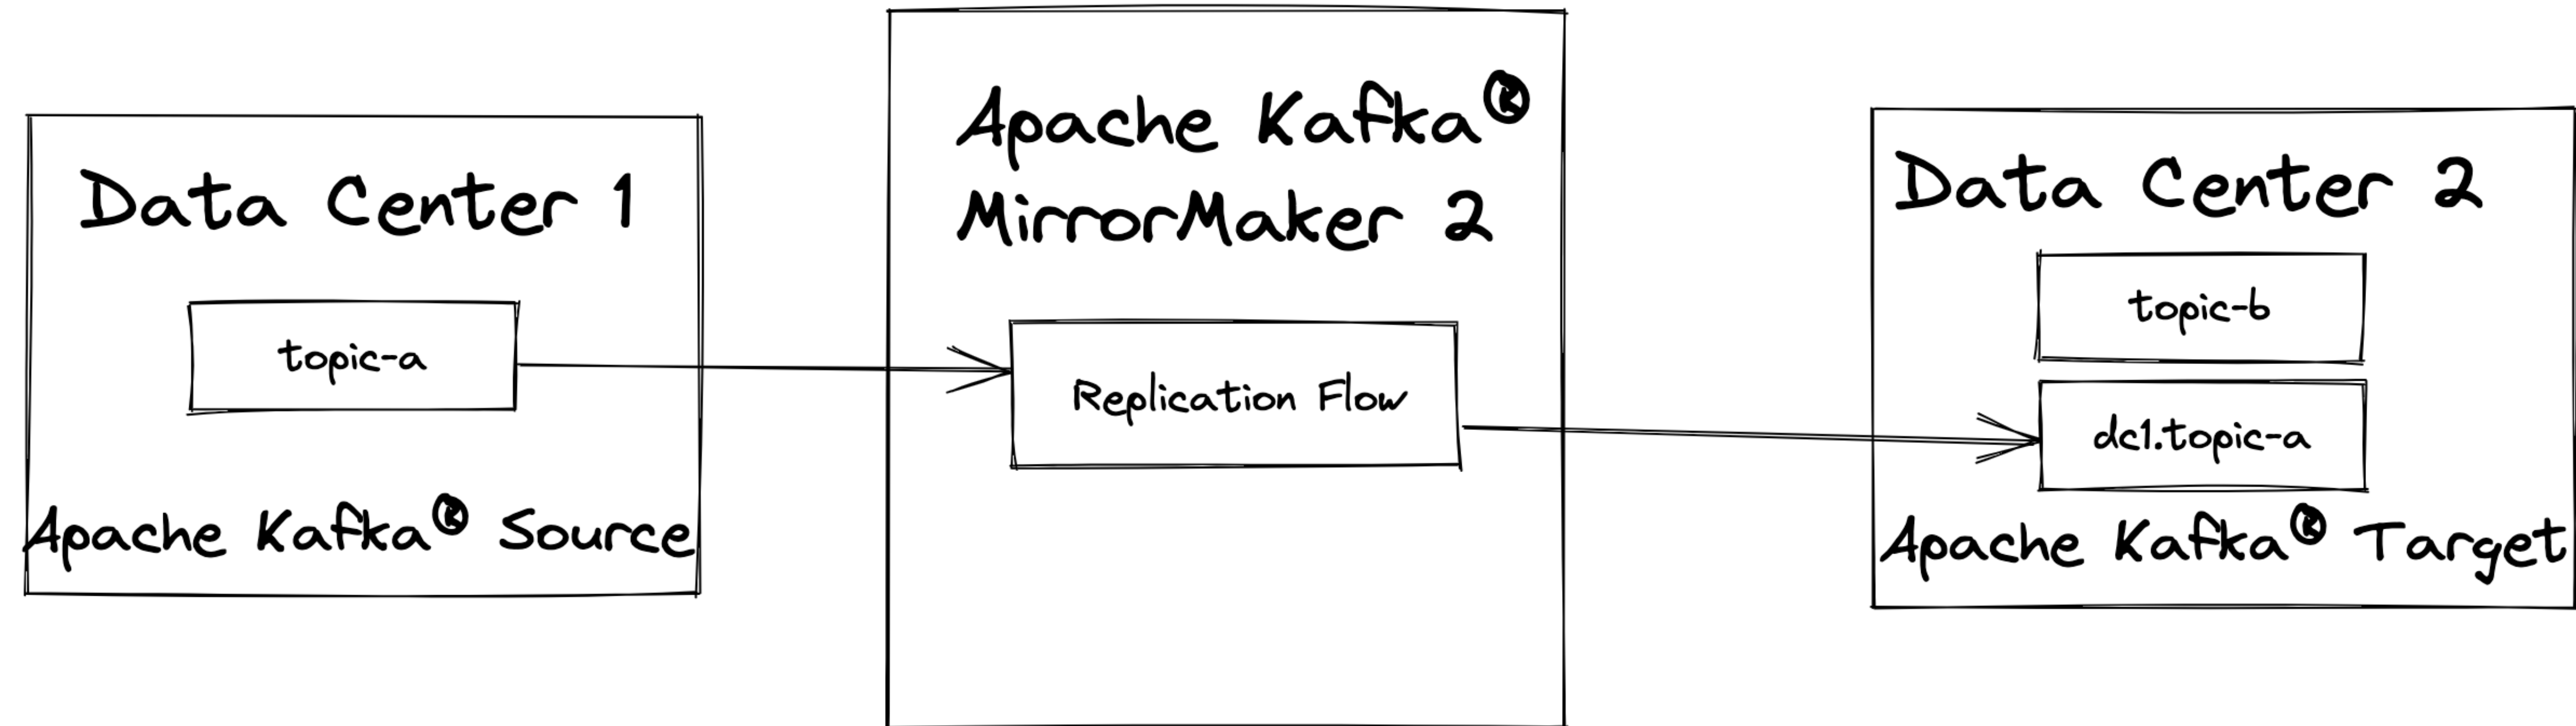
Repeatability

Consistency

@aiven_io

in/diahmed | @DewanAhmed

Apache Kafka® and Apache Kafka® MirrorMaker 2



There was a demo in this slide.

<https://developer.aiven.io/docs/tools/terraform/reference/cookbook/kafka-mirrormaker-recipe.html> to follow the demo

But... How do I start?

- Research - is this right for us?
- Awareness - are we ready for this?
- Principles - Do you know the fundamentals?
- Pilot - Let's get something working
- ~~All-in~~ Strategy - Roll-out where it makes sense

Recap

- Infrastructure-as-Code (IaC) principles
- Challenges and concerns - IaC with data infrastructure
- How to get started in your IaC journey?

References:

- Aiven Terraform Provider: <https://developer.aiven.io/docs/tools/terraform.html>
- Demo: <https://developer.aiven.io/docs/tools/terraform/reference/cookbook/kafka-mirrormaker-recipe.html>
- IaC Blog: <https://www.dewanahmed.com/iac-principles-patterns>

Questions?

dewan@aiven.io